

# Untitled

December 10, 2025

```
[1]: #Importing the necessary libraries
```

```
[2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from datetime import datetime
from sklearn.cluster import KMeans
from sklearn.decomposition import PCA
```

```
[3]: df1 = pd.read_csv('dailyActivity_merged.csv')
print(df1)
```

	Id	ActivityDate	TotalSteps	TotalDistance	TrackerDistance	\
0	1503960366	4/12/2016	13162	8.500000	8.500000	
1	1503960366	4/13/2016	10735	6.970000	6.970000	
2	1503960366	4/14/2016	10460	6.740000	6.740000	
3	1503960366	4/15/2016	9762	6.280000	6.280000	
4	1503960366	4/16/2016	12669	8.160000	8.160000	
..	..	..	..	..	..	
935	8877689391	5/8/2016	10686	8.110000	8.110000	
936	8877689391	5/9/2016	20226	18.250000	18.250000	
937	8877689391	5/10/2016	10733	8.150000	8.150000	
938	8877689391	5/11/2016	21420	19.559999	19.559999	
939	8877689391	5/12/2016	8064	6.120000	6.120000	
	LoggedActivitiesDistance	VeryActiveDistance	ModeratelyActiveDistance			\
0	0.0	1.88		0.55		
1	0.0	1.57		0.69		
2	0.0	2.44		0.40		
3	0.0	2.14		1.26		
4	0.0	2.71		0.41		
..	..	..	..	..	..	
935	0.0	1.08		0.20		
936	0.0	11.10		0.80		
937	0.0	1.35		0.46		
938	0.0	13.22		0.41		
939	0.0	1.82		0.04		

	LightActiveDistance	SedentaryActiveDistance	VeryActiveMinutes	\
0	6.06	0.00	25	
1	4.71	0.00	21	
2	3.91	0.00	30	
3	2.83	0.00	29	
4	5.04	0.00	36	
..	...	...	...	
935	6.80	0.00	17	
936	6.24	0.05	73	
937	6.28	0.00	18	
938	5.89	0.00	88	
939	4.25	0.00	23	
	FairlyActiveMinutes	LightlyActiveMinutes	SedentaryMinutes	Calories
0	13	328	728	1985
1	19	217	776	1797
2	11	181	1218	1776
3	34	209	726	1745
4	10	221	773	1863
..	...	...	...	...
935	4	245	1174	2847
936	19	217	1131	3710
937	11	224	1187	2832
938	12	213	1127	3832
939	1	137	770	1849

[940 rows x 15 columns]

```
[4]: df1["Date"] = pd.to_datetime(df1["ActivityDate"])
df1 = df1.drop(columns=["ActivityDate"])

[5]: df1 = df1[["Id", "Date", "TotalSteps", "Calories", "TotalDistance",
   ↪ "VeryActiveMinutes", "FairlyActiveMinutes",
   ↪ "LightlyActiveMinutes", "SedentaryMinutes"]]

df1 = df1.drop_duplicates()
df1.duplicated().sum()
df1.isnull().sum()
print(df1.head())
print(df1.shape)
print(df1.columns)
#dropping the duplicates from the data (the same way it was taught at the
#course assignments) and checking that there are only lines with information
#included.
>Selecting the columns that I want to use in the project
```

	Id	Date	TotalSteps	Calories	TotalDistance	\
--	----	------	------------	----------	---------------	---

```

0 1503960366 2016-04-12      13162      1985      8.50
1 1503960366 2016-04-13      10735      1797      6.97
2 1503960366 2016-04-14      10460      1776      6.74
3 1503960366 2016-04-15      9762       1745      6.28
4 1503960366 2016-04-16     12669      1863      8.16

    VeryActiveMinutes  FairlyActiveMinutes  LightlyActiveMinutes  \
0                  25                  13                  328
1                  21                  19                  217
2                  30                  11                  181
3                  29                  34                  209
4                  36                  10                  221

   SedentaryMinutes
0              728
1              776
2             1218
3              726
4              773
(940, 9)
Index(['Id', 'Date', 'TotalSteps', 'Calories', 'TotalDistance',
       'VeryActiveMinutes', 'FairlyActiveMinutes', 'LightlyActiveMinutes',
       'SedentaryMinutes'],
      dtype='object')

```

[6]: `#this checks that we are using the data from the 31 days between 12.4-12.5.2016`

```

[7]: df1 = df1[
        (df1["Date"] >= "2016-04-12") &
        (df1["Date"] <= "2016-05-12")
    ]
print(df1.shape)
print(df1["Date"].nunique())
print(df1["Date"].min(), df1["Date"].max())

```

```

(940, 9)
31
2016-04-12 00:00:00 2016-05-12 00:00:00

```

[8]: `print(df1) #here we have the data that is cleaned and ready for visualizations`

```

          Id      Date  TotalSteps  Calories  TotalDistance  \
0  1503960366 2016-04-12      13162      1985      8.500000
1  1503960366 2016-04-13      10735      1797      6.970000
2  1503960366 2016-04-14      10460      1776      6.740000
3  1503960366 2016-04-15      9762       1745      6.280000
4  1503960366 2016-04-16     12669      1863      8.160000
..      ...

```

935	8877689391	2016-05-08	10686	2847	8.110000
936	8877689391	2016-05-09	20226	3710	18.250000
937	8877689391	2016-05-10	10733	2832	8.150000
938	8877689391	2016-05-11	21420	3832	19.559999
939	8877689391	2016-05-12	8064	1849	6.120000
					\
0		VeryActiveMinutes	FairlyActiveMinutes	LightlyActiveMinutes	
1		25	13	328	
2		21	19	217	
3		30	11	181	
4		29	34	209	
..		..	..	..	
935		36	10	221	
936		17	4	245	
937		73	19	217	
938		18	11	224	
939		88	12	213	
		23	1	137	
		SedentaryMinutes			
0		728			
1		776			
2		1218			
3		726			
4		773			
..		..			
935		1174			
936		1131			
937		1187			
938		1127			
939		770			

[940 rows x 9 columns]

```
[9]: df5 = pd.read_csv('sleepDay_merged.csv')
df5["Date"] = pd.to_datetime(df5["SleepDay"])
print(df5)
#this reads the sleepdata
```

	Id	SleepDay	TotalSleepRecords	TotalMinutesAsleep	\
0	1503960366	4/12/2016 12:00:00 AM	1	327	
1	1503960366	4/13/2016 12:00:00 AM	2	384	
2	1503960366	4/15/2016 12:00:00 AM	1	412	
3	1503960366	4/16/2016 12:00:00 AM	2	340	
4	1503960366	4/17/2016 12:00:00 AM	1	700	
..	..	..	..	..	
408	8792009665	4/30/2016 12:00:00 AM	1	343	
409	8792009665	5/1/2016 12:00:00 AM	1	503	

```

410 8792009665 5/2/2016 12:00:00 AM 1 415
411 8792009665 5/3/2016 12:00:00 AM 1 516
412 8792009665 5/4/2016 12:00:00 AM 1 439

```

	TotalTimeInBed	Date
0	346	2016-04-12
1	407	2016-04-13
2	442	2016-04-15
3	367	2016-04-16
4	712	2016-04-17
..	..	..
408	360	2016-04-30
409	527	2016-05-01
410	423	2016-05-02
411	545	2016-05-03
412	463	2016-05-04

[413 rows x 6 columns]

/tmp/ipykernel\_654/3438589101.py:2: UserWarning: Could not infer format, so each element will be parsed individually, falling back to `dateutil`. To ensure parsing is consistent and as-expected, please specify a format.

```
df5["Date"] = pd.to_datetime(df5["SleepDay"])
```

```
[10]: df5 = df5[
    (df5["Date"] >= "2016-04-12") &(df5["Date"] <= "2016-05-12")]
print(df5.shape)
print(df5["Date"].min(), df5["Date"].max())
#this checks that the sleepdata is also from 12.4-12.5. period
```

(413, 6)

2016-04-12 00:00:00 2016-05-12 00:00:00

```
[11]: df5 = df5[[ "Id", "Date", "TotalMinutesAsleep", "TotalTimeInBed"]]
df5.duplicated().sum()
df5 = df5.drop_duplicates()
#dropping duplicates from the sleepdata
```

```
[12]: df = pd.merge(df1, df5, on=[ "Id", "Date"], how="left")
print(df) # print the compined data
```

	Id	Date	TotalSteps	Calories	TotalDistance	\
0	1503960366	2016-04-12	13162	1985	8.500000	
1	1503960366	2016-04-13	10735	1797	6.970000	
2	1503960366	2016-04-14	10460	1776	6.740000	
3	1503960366	2016-04-15	9762	1745	6.280000	
4	1503960366	2016-04-16	12669	1863	8.160000	
..	..	..	..	..	..	

935	8877689391	2016-05-08	10686	2847	8.110000
936	8877689391	2016-05-09	20226	3710	18.250000
937	8877689391	2016-05-10	10733	2832	8.150000
938	8877689391	2016-05-11	21420	3832	19.559999
939	8877689391	2016-05-12	8064	1849	6.120000

	VeryActiveMinutes	FairlyActiveMinutes	LightlyActiveMinutes	\
0	25	13	328	
1	21	19	217	
2	30	11	181	
3	29	34	209	
4	36	10	221	
..	...	...	...	
935	17	4	245	
936	73	19	217	
937	18	11	224	
938	88	12	213	
939	23	1	137	

	SedentaryMinutes	TotalMinutesAsleep	TotalTimeInBed
0	728	327.0	346.0
1	776	384.0	407.0
2	1218	NaN	NaN
3	726	412.0	442.0
4	773	340.0	367.0
..	...	...	...
935	1174	NaN	NaN
936	1131	NaN	NaN
937	1187	NaN	NaN
938	1127	NaN	NaN
939	770	NaN	NaN

[940 rows x 11 columns]

```
[13]: # first visualization: boxplot of the average steps compared to WHO
      ↴recommendations
average_steps = df1['TotalSteps'].mean()
comparison_data = {
    "Category": ["Fitbit group average", "WHO recommendation"],
    "Steps": [average_steps, 10000]
}
comparison_df = pd.DataFrame(comparison_data)
plt.figure(figsize=(6,5))
ax = sns.barplot(data=comparison_df, x="Category", y="Steps", color = "green")


for p in ax.patches: #adds the numbers of steps on top of the bar plots
    ax.annotate(
        str(p.get_height()),
        (p.get_x() + p.get_width() / 2, p.get_height()),
        ha='center',
        va='bottom'
    )
```

```

    format(int(p.get_height())),
    (p.get_x() + p.get_width() / 2., p.get_height()), color= "blue")

plt.title("Average steps vs. WHO recommendation")
plt.ylabel("Steps per day")
plt.xlabel("")
plt.show()

```



```

[14]: avg_daily_steps = df1.groupby("Date")["TotalSteps"].mean().reset_index()
plt.figure(figsize=(10,6))

# creating a lineplot of average daily steps compared to the 10 000 steps
sns.lineplot(
    data=avg_daily_steps,
    x="Date",
    y="TotalSteps",
    marker="o",
    color = "red",
    label = "Average daily steps of FitBit trackers",
)

```


```

#this draws the green line for the WHO recommendation
plt.axhline(y=10000, linestyle="--", color="green", label="WHO recommendation\u219210,000 steps / per day")

plt.title("Average daily steps vs. WHO recommendation", fontsize=14)
plt.xlabel("Date")
plt.ylabel("Average steps")
plt.xticks(rotation=45) #this makes the text rotate on the x-axis to be easier to read
plt.legend()
plt.grid(True, linestyle="--", alpha=0.5)

plt.tight_layout()
plt.show()

```



[15]: df\_sleep = df.dropna(subset=["TotalMinutesAsleep"])


[16]: plt.figure(figsize=(8,6))
sns.scatterplot(data=df\_sleep,
 x="TotalSteps",
 y="TotalMinutesAsleep")

plt.title("Total steps vs. sleep duration")
plt.xlabel("Total steps")

```

plt.ylabel("Minutes asleep")
plt.tight_layout()
plt.show()
#extra visualisation not used in the report

```



```

[17]: df['TotalActiveMinutes'] = (df['VeryActiveMinutes'] + df['FairlyActiveMinutes']
                                + df['LightlyActiveMinutes'])
corr = df["TotalActiveMinutes"].corr(df["Calories"])
#combining all the active minutes to create one variable TotalActiveMinutes and
#then calculate the correlation between the two columns

plt.figure(figsize=(8,6))
sns.regplot(
    data=df,
    x="TotalActiveMinutes",
    y="Calories",
    scatter_kws={"alpha":0.5}, #this makes the blue dots to be a little
    #see through in order to see dots better when they are overlapping
    line_kws={"color": "red"}
)


```

```

plt.text(
    x=50,
    y=df["Calories"].max() - 200,
    s=f"Correlation: {corr:.2f}", #corr:.2f shows correlation with 2 decimals
    fontsize=12,
    color="black"
)

plt.title("Total active minutes vs. calories burned")
plt.xlabel("Total active inutes")
plt.ylabel("Calories")
plt.tight_layout()
plt.show()

```



```

[18]: df_sleep["Weekday"] = df_sleep["Date"].dt.day_name()

plt.figure(figsize=(8,6))
#creating a distribution boxplot of sleep duration across week days, boxplot_
    ↴was introduced at the course material
sns.boxplot(data=df_sleep,

```

```


x="Weekday",
y="TotalMinutesAsleep",
order=["Monday", "Tuesday", "Wednesday", "Thursday", "Friday", "Saturday", "Sunday"], color = "green")

plt.title("Distribution of sleep duration across week days")
plt.xlabel("Day")
plt.ylabel("Minutes asleep")
plt.xticks(rotation=30)
plt.tight_layout()
plt.show()

```

/tmp/ipykernel\_654/1340094653.py:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead


See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)  
df\_sleep["Weekday"] = df\_sleep["Date"].dt.day\_name()  
/opt/software/lib/python3.10/site-packages/seaborn/categorical.py:632:  
FutureWarning: SeriesGroupBy.grouper is deprecated and will be removed in a  
future version of pandas.  
positions = grouped.grouper.result\_index.to\_numpy(dtype=float)



```
[19]: df_sleep = df_sleep.copy()
df_sleep.loc[:, "TimeInBedHours"] = df_sleep["TotalTimeInBed"] / 60 #change time to hours

sns.lmplot(data=df_sleep, x="TimeInBedHours", y="TotalMinutesAsleep", line_kws={"color": "green"})

plt.xlabel("Time in bed (hours)")
plt.ylabel("Minutes asleep")
plt.title("Time in bed vs. sleep duration")
plt.show()
```



```
[20]: from sklearn.preprocessing import StandardScaler
#advanced analysis method the k-means clustering, it was introduced at the
#course assignment material
clustering_data = df[["TotalSteps", "TotalMinutesAsleep"]].dropna() #first we
#drop the null values from the columns we are using

scaler = StandardScaler() #using standard scaler to scale the different
#measurement types
scaled = scaler.fit_transform(clustering_data)

kmeans = KMeans(n_clusters=3, random_state=42) #3 defines that we want 3
#different clusters
clusters = kmeans.fit_predict(scaled)


clustering_data["Cluster"] = clusters
```

```

plt.scatter(clustering_data["TotalSteps"],
            clustering_data["TotalMinutesAsleep"],
            c=clustering_data["Cluster"])
plt.xlabel("Steps")
plt.ylabel("Sleep minutes")
plt.title("Clusters based on activity and sleep")
plt.show()

```

/opt/software/lib/python3.10/site-packages/sklearn/cluster/\_kmeans.py:1416:  
 FutureWarning: The default value of `n\_init` will change from 10 to 'auto' in  
 1.4. Set the value of `n\_init` explicitly to suppress the warning  
 super().\_\_check\_params\_vs\_input(X, default\_n\_init=10)



```

[21]: corr_cols = ["TotalSteps", "Calories", "TotalMinutesAsleep", "TotalTimeInBed",  

   ↴ "SedentaryMinutes", "VeryActiveMinutes"]  

#corr_cols defines all the columns I want to use in the matrix  

corr_df = df[corr_cols].dropna()  

#creating the correlation heatmap, corr was introduced in the assignment  


   ↴ materials

```

```

sns.heatmap(corr_df.corr(), annot=True, fmt=".2f", cmap="Greens") #fmt=".2f" ↴
    ↴defines that the matrix values are only presented with two decimals
plt.title("Correlation heatmap of activity and sleep variables")
plt.show()

```



```

[22]: df_min = pd.read_csv("minuteStepsNarrow_merged.csv")

df_min["ActivityMinute"] = pd.to_datetime(
    df_min["ActivityMinute"],
    format="%m/%d/%Y %I:%M:%S %p" #transforming the time to correct format
)

df_min["Hour"] = df_min["ActivityMinute"].dt.hour

# calculating the average for hour
hourly_steps = df_min.groupby("Hour")["Steps"].mean().reset_index()


```

```

plt.figure(figsize=(10,6))
sns.lineplot(data=hourly_steps, x="Hour", y="Steps", marker="o")

plt.title("Average steps per hour of day ")
plt.xlabel("Hour of day")
plt.ylabel("Average steps")
plt.xticks(range(0,24))
plt.grid(True, linestyle="--", alpha=0.4)
plt.show()

```



```


[23]: user_id = 1503960366 #just random user id from the data
df_subject = df[df['Id'] == user_id]
plt.figure(figsize=(8,6))

#individual level visualisation
sns.lineplot(data=df_subject, x = "Date", y = "TotalSteps", marker = "o")

plt.title(f"User 1503960366 daily step count over the observation period")
plt.xlabel("Date")
plt.ylabel("Total steps")
plt.xticks(rotation=45) #this rotates the text in x-axis so its easier to read
plt.tight_layout()
plt.show()

```

User 1503960366 daily step count over the observation period



[ ]:

[ ]:

[ ]: